

DOCUMENTATION FOR THE RATPOINTS PROGRAM

MICHAEL STOLL

1 Introduction

This paper describes the `ratpoints` program. This program tries to find all rational points within a given height bound on a hyperelliptic curve in the most efficient way possible.

2 History and Acknowledgments

This program goes back to an implementation of the ‘quadratic sieving’ idea by **Noam Elkies** that was around in the early 1990s. My own first contribution was to replace the char arrays that were used to store the sieving information by bit arrays, in 1995. **Colin Stahlke** then made use of the `gmp` library, so that points could be checked exactly, and implemented the selection of sieving primes according to their likely success rate, in 1998. After that, I successively put in numerous improvements (and some bug fixes). For details of how the program works, see Section 7 below.

Along with **Noam Elkies** and **Colin Stahlke**, I would like to thank **John Cremona** and **Sophie Labour** for bug reports and suggestions for improvements. Further thanks are due to **Bill Allombert** for a first implementation of the use of 256-bit (and possibly 512-bit) registers.

3 Availability

The `ratpoints-3.0.0` package can be downloaded from my homepage, see [rat], or obtained from github at <https://github.com/MichaelStollBayreuth/ratpoints>.

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but **without any warranty**; without even the implied warranty of **merchantability** or **fitness for a particular purpose**. See the GNU General Public License for more details.

You should receive a copy of the GNU General Public License along with this program.

If not, see <http://www.gnu.org/licenses/>.

4 Installation

This section describes the installation procedure under Linux.

4.1 Extract the archive

```
> tar xzf ratpoints-3.0.0.tar.gz
```

This sets up a directory `ratpoints-3.0.0` containing the various files that belong to the installation.

Date: September 19, 2026.

4.2 Build the program and library

Do

```
> cd ratpoints-3.0.0
> make all
```

This will build the library `libratpoints.a`, the executable `ratpoints`, and also run `pdflatex` to generate this documentation.

Five of the constants that govern the sieve depend on the machine rather than on the curve. Reasonable values are compiled in, but `make tune` will measure better ones for your machine if you want it to; see Section 8.4. It is a separate step because it takes several minutes.

By default, the program has the odd primes below 256 available and considers the smallest 30 of them; for a curve that needs more, it will look further of its own accord, as described in Section 8.1. If you intend to do computations with large height bounds, it may make sense to make more primes available. This can be achieved via

```
> make all PRIME_SIZE=9
```

The `PRIME_SIZE` argument can be given any value from 5 to 10; otherwise it is taken to be 8. The precise meaning is that the program will work with primes $< 2^s$, when `PRIME_SIZE = s`. The `Makefile` notices that the setting has changed and rebuilds what depends on it. Only the primes a run actually asks for get sieving tables, so a larger table costs nothing until it is used (Section 5.5).

The sieve works on several numerators at a time, in a register whose width is selected by the line

```
CCFLAGS1 = ${CCFLAGS256}
```

in `Makefile`. It can be set to `${CCFLAGS64}` (64-bit registers, i.e., plain unsigned longs, which works on any machine the library builds on at all), `${CCFLAGS128}`, `${CCFLAGS256}` (the default, using 256-bit AVX registers) or `${CCFLAGS512}`. The `Makefile` notices that the setting has changed and rebuilds everything that depends on it, so a `make distclean` in between is not necessary. Wider registers help the first phase of the sieve, where most of the time goes on a curve with few rational points, but they help the second phase much less and the sieve tables not at all, so the gain is nowhere near proportional to the width: on the author's machine, 64, 128 and 256 bits take 21.3, 15.2 and 12.2 units of time on such a curve.

The 128-bit setting uses the macro `USE_AVX128`. In spite of its name it needs only SSE2 instructions, so it runs wherever the older `USE_SSE` variant did, and it is never slower, because its test for an empty register compares the whole register against zero instead of extracting both halves into general registers. That is worth about 30% of the second phase of the sieve, which comes to 9% of a run long enough for the sieve to dominate it and about 3% of `make test1`. The older variant is still available as `${CCFLAGS128s}`.

You may have to change `-mavx2` in the line

```
CCFLAGS256 = -DUSE_AVX -mavx2
```

into `-mavx` or leave it out altogether, depending on the capabilities of your CPU. You will notice that AVX2 (or AVX) is not supported, when the executable throws an 'unknown instruction' error or similar.

There is also a variant using 512-bit registers, activated by

```
CCFLAGS1 = ${CCFLAGS512}
```

which requires a CPU with AVX512F capability. Note that this variant has so far only been tested indirectly (see the change log entry below), since I have no such CPU available; you should therefore run `make test` before relying on it. It is also worth timing the 256-bit and the 512-bit version against each other on your machine: the sieving loop is limited by how fast it loads bit arrays from the first-level cache and by how much of that cache the sieve tables occupy, rather than by arithmetic, so wider registers do not automatically pay off, and some Intel server processors reduce their clock frequency when 512-bit instructions are used. The one data point available so far is encouraging, though: on an AMD Ryzen 9 9950X3D (Zen 5 architecture), **Drew Sutherland** measured a speed-up of 13 to 14% over the 256-bit version, for a curve with many rational points. If you leave out `-mavx512f` and compile with just `-DUSE_AVX512`, then `gcc` will express the 64-byte vectors through narrower operations, so that the code runs on any machine. This is of no use in production (it is slower than the 256-bit version), but it does allow the 512-bit code path to be tested without the corresponding hardware. `gcc` will warn that passing a 512-bit vector argument without AVX512F enabled changes the ABI; this is harmless, as the whole library is built with the same flags.

4.3 Run various tests

Run

```
> make test1
```

in the working directory. This will build an executable `rptest` and then run (and time) it. Finally, the output (which was written to a file `rptest.out`) is compared against `testbase`, which contains the output of a sample run. The two should be identical; otherwise the message `Test failed!` will be displayed (on its own on a line) and the target fails, so that `make` exits with an error and names the test.

```
> make test1once
```

runs the same curves again, but with the input fields of the argument structure set once before the loop over the curves instead of once per curve, as a program using the library may do since version 3.0.0 (Section 6); `rptest` prints a line whenever the library has changed one of those fields, and the output is compared against `testbase` as before.

```
> make test1many
```

does the same with a second collection of curves, in `testdata-many.h`, whose reference output is `testbase-many`. These curves have (comparatively) many rational points, which makes them sieve quite differently from the random curves in `testdata.h`: on a random curve about one numerator in 10^4 survives the first sieving stage, on these up to a hundred times more, so the best division of labour between the two stages is a different one. The two tests take about the same time and between them cover the two regimes that occur in practice; both should be used whenever the parameters are retuned, since a setting that suits one of them can be a poor choice for the other. That is what `make tune` does with them; see Section 8.

```
> make testhigh
```

```
> make testhighmany
```

run the same two regimes at a height bound of $2 \cdot 10^5$ rather than the 16383 of `test1` and `test1many`, and take about a minute each. That bound decides a good deal of what a test measures. A sieve table is built the first time a given prime meets a given class of denominators and is reused afterwards, so the total cost of building the tables is bounded by the number of primes and does not grow with the height, whereas the sieving does: on the random curves the sieve tables are 4%

of make `test1` (they were 22% when these suites were added, before the set-up was reworked) but 0.3% of make `testhigh`, the whole of the set-up per denominator 8.5% against 1.3%, and the two sieving phases themselves go from 57% to 90% of the run. These are therefore the tests to judge a change to the sieving loops by. `testhigh` uses the same curves as `test1`; none of them has a rational point of height between 16383 and $2 \cdot 10^5$, so the expected output is `testbase` again, and that the two agree is itself worth knowing. `testhighmany` uses the thirty curves of `testdata-many.h` that run out of sieving primes, collected in `testdata-high-many.h` with reference output `testbase-high-many`.

> make `test2`

calls `ratpoints` on a curve with lots of rational points and with a fairly large height bound, times it, and compares the output to what is expected. This may take a few minutes, so be patient!

> make `test3`

runs the invocations of `ratpoints` collected in `test3.sh`, one or more for each of the bugs fixed in versions 2.2.4 and 3.0.0, and compares the output against `testbase3`.

> make `test4`

runs the invocations of `ratpoints` collected in `test4.sh`, a few hundred of them, chosen so that every branch of the code runs at least once: degrees 1 to 9 and up to the maximum of 100, odd and even, square and non-square leading and constant coefficients, every reason to reverse the polynomial, curves with no points modulo some prime or no admissible numerator modulo 64, forbidden divisors of every kind, every option, restricted denominator ranges and search intervals, height bounds from 1 to $2^{63} - 1$ (those come with a denominator range of one and a narrow search interval, so that the run stays short), and every error message. The output is compared against `testbase4`; the whole test takes about a second. Before each invocation the script prints its arguments, so that a failing test can be found from the difference between `test4.out` and `testbase4`. The reference was checked independently of the sieve: `verify-test4.py` reads it and, for every invocation that prints points, searches every coprime pair (a, b) in the range by brute force, in arbitrary precision; it can be run again whenever the reference changes (it needs Python 3, and PARI/GP for the 19 searches, on two curves, that are too large for it).

> make `testapi`

runs `rpapi`, a program that tests what the library offers and the command-line program cannot reach: the checks that `find_points_work` makes on its arguments and the error codes it returns, a callback that stops the search or declines a point, the flags, and a sequence of searches on one initialized structure (Section 6); the output is compared against `testbase-api`.

> make `test4configs`

runs `test4` again on the library built with each of the other compile-time switches: the register widths 64, 128 (both variants) and 512 (emulated); `RATPOINTS_CHUNK=1`, one register in the first phase; `USE_LONG_IN_PHASE_2`, the second phase one word at a time; `PRIME_SIZE 7`, the smaller prime table; and composite moduli off and unbounded. Each is built in a directory of its own, so that the build in the working directory is left alone (`test4-configs.sh`); the points found must not depend on any of these, so every build is compared against the same `testbase4`; the build with the smaller table is compared without the reports of `-v`, two of which say what the primes beyond 127 say. About a minute. The script exits with status 1 when an output differs from the reference and with 2 when a build failed; so does `coverage.sh` below.

> make `coverage`

measures how much of the code the tests execute, with `gcov`: it builds the sources instrumented and unoptimised in `build-coverage/`, runs the tests through that build (all of `make test` but `test2` and `timing`, which add nothing but minutes) and prints, for each source file, the lines executed and the branches taken at least once (`coverage.sh`); the annotated sources are left in that directory. When this was written the tests executed 99.9% of the lines of `find_points.c`, 96% of `sift.c`, all of `init.c` and `sturm.c` and 99% of `main.c`; what is left is guards against states no caller can produce, arms that other compile-time settings reach (`test4configs` runs those builds, but this measures one), and a few conditions the arithmetic makes impossible.

```
> make timing
```

times `ratpoints` on some curve with height bound 400 000. This maybe useful for comparisons.

Finally,

```
> make test
```

runs `test1`, `test1once`, `test1many`, `testdegrees`, `test2`, `test3`, `test4`, `testapi` and `timing`, each of them whatever the earlier ones did, and fails at the end if any of them failed, so that a script can tell whether the build passed. The two large-height tests are not part of it; they are meant for judging a change to the sieve rather than for checking that the build works.

4.4 Install

If you like, you can install the library, executable and header file on your system.

```
> sudo make install
```

The executable is copied to `/usr/local/bin/`, the library to `/usr/local/lib/`, and the header file to `/usr/local/include/`. You can change the `/usr/local` prefix via

```
> make install INSTALL_DIR=...
```

4.5 Debugging

In case you found a bug and would like to find out where it comes from, or if you just want to see exactly what the program is doing, you can do

```
> make debug
```

in the working directory. This will build an executable `ratpoints-debug`, which, when run, will dump loads of output on the screen, so it is best to send the output to a file, which you can then study at leisure.

4.6 Cleaning up

In order to get rid of the temporary files, do

```
> make clean
```

To remove everything except the files from the archive, use

```
> make distclean
```

4.7 Requirements

You need a C compiler (like `gcc`, which is the default specified for the `CC` variable in `Makefile`; if necessary, it can be changed there). Since version 3.0.0, the code assumes that `long` has 64 bits and refuses to compile otherwise; that is the case on every 64-bit Unix-like system (but not under the Windows ABI, where `long` has 32 bits even on a 64-bit machine). Version 2.2.4 is the last one that accommodates a 32-bit `long`.

In addition to the standard libraries, the program also requires the `gmp` (GNU multi-precision) library [\[gmp\]](#).

4.8 List of files

The archive contains the following files.

- `Makefile`
- `ratpoints.h` — the header file for programs using `ratpoints`
- `rp-private.h`, `primes.h` — header files used internally
- `gen_find_points.h.c`, `gen_init_sieve.h.c` — short programs that write additional header files depending on the system configuration
- `sift.c`, `init.c`, `sturm.c`, `find_points.c` — the source code for the `ratpoints` library
- `main.c`, `rpctest.c`, `rpapi.c` — the source code for the `ratpoints`, `rpctest` and `rpapi` executables, respectively
- `testdata.h`, `testdata-many.h`, `testdata-high-many.h`, `testdata-degrees.h` — the curves used by the test runs
- `testbase`, `testbase-many`, `testbase-high-many`, `testbase-degrees`, `testbase2`, `testbase3`, `testbase4`, `testbase-api` — the output they are expected to produce
- `test3.sh`, `test4.sh` — the scripts that make `test3` and `test4` run; they have to be executable, as have the next three
- `verify-test4.py` — the brute-force check of `testbase4` (needs Python 3)
- `test4-configs.sh`, `coverage.sh` — what make `test4configs` and make `coverage` run
- `tune.sh` — what make `tune` runs to measure the machine-dependent constants (Section 8)
- `bench_init.c`, `bench_check.c` — benchmarks for two parts that are timed on their own: building the sieving tables, and the exact test. Neither is part of the library; the second is how to remeasure the constants of Section 8.2
- `ratpoints-doc-3.0.pdf` — this documentation file
- `gpl-2.0.txt` — the GNU license that applies to this program.

5 How to use `ratpoints`

5.1 Basic operation

Let

$$C : y^2 = a_n x^n + a_{n-1} x^{n-1} + \cdots + a_1 x + a_0$$

be your curve, and let H be the bound for the denominator and absolute value of the numerator of the x -coordinate of the points you want to find. The command to search for the points is then

```
> ratpoints 'a_0 a_1 ... a_{n-1} a_n' H
```

The first argument to `ratpoints` is the list of the coefficients, which have to be integers, are separated by spaces, and are listed starting with the constant term. There is no bound on the size of the coefficients; they are read as multi-precision integers.

The degree n of the polynomial is limited to `RATPOINTS_MAX_DEGREE` (defined in `ratpoints.h`), which is set to 100 by default.

The program will give an error message when the polynomial (considered as a binary form of the smallest even degree $\geq n$) is not squarefree (e.g., if you specify two leading zero coefficients).

The following subsections discuss how to modify the standard behavior. This is achieved by adding options after the two required arguments. Some of the options have arguments, others don't; the ordering of the options and option-argument pairs is largely arbitrary. The exceptions are options with opposite effect, where the last one given counts, and the specification of the search intervals, which must be given in order.

5.2 Early abort

By default, `ratpoints` will search the whole region that you specify and print all the points it finds. If you just want to find *one* point (for example when dealing with 2-covering curves of elliptic curves), you can tell the program to quit after it has found one point via the `-1` option, e.g.,

```
> ratpoints '-18 116 48 -12 30' 60000000 -1
```

If you don't want to see points at infinity, specify the `-i` option. The two options can be combined: `-i -1` will stop the program after it has found one finite point.

5.3 Changing the output

By default, `ratpoints` prints some general information before and after the points, which are given in the form $(x : y : z)$, one per line. Here, $(x : y : z)$ are the coordinates of the point considered as a point in the $(1, \lceil n/2 \rceil, 1)$ -weighted projective plane, which is the natural ambient space for the curve. The coordinates are integers with x and z coprime and $z > 0$, or $z = 0$ and $x = 1$ (for points at infinity).

There are various ways to change this behavior.

- To suppress all output except the points, use `-q` (for *quiet*).
- To add some more output explaining what the program is doing, use `-v` (for *verbose*). This has no effect if the `-q` option is also given.
- To suppress printing of the points, use `-z`.
- If you only want to list the x -coordinates of point pairs rather than individual points, use `-y`.
- There are four options that influence the format in which the points are printed:
`-f format -fs string-before -fm string-between -fe string-after`

The arguments to all of them are strings; `"\n"`, `"\t"`, `"\\"` and `"\%"` are recognized and do what you expect. The *format* string can contain markers `%x`, `%y`, `%z` that will be replaced with the x , y , and z -coordinate of the point, respectively. The defaults are empty for *string-before*, *string-between* and *string-after*, and `"(%x : %y : %z)\n"` for *format* when `-y` is not specified; otherwise `"(%x : %z)\n"`.

The effect is as follows. Before any point is printed, *string-before* is output. Then every point is printed according to *format*. Between any two points, *string-between* is output, and after the last point, *string-after* is output.

As an example, consider the effect of

```
-f "[%x,%y,%z]" -fs "{" -fm "," -fe "}"\n
```

5.4 Restricting the search domain

There are two ways to restrict the domain of the search. The first is to restrict the range of denominators considered. This is done via

`-dl dmin -du dmax`.

For example, to only look for integral points, you can say `-du 1`. By default, the lower limit is 1 and the upper limit is H .

The other way is to restrict the search to a union of (closed) intervals. If you only want to search for points in

$$[l_1, u_1] \cup [l_2, u_2] \cup \cdots \cup [l_k, u_k],$$

you can specify this in the form

`-l l1 -u u1 -l l2 -u u2 ... -l lk -u uk`.

The first `-l` option is optional; l_1 defaults to $-\infty$. Similarly the last `-u` option is optional, and u_k defaults to $+\infty$. The number k of intervals is bounded by `RATPOINTS_MAX_DEGREE` (usually, 100).

5.5 Setting the parameters for the sieve

It is possible to change the number of primes that are used in the various stages of the algorithm. This is done by the following options.

`-p M -N N -n n`

This sets the number of (odd) primes that are considered for the sieve to M , the number of primes that are actually used for the sieve to N , and the number of primes used in the first sieving stage to n . Giving M also fixes it: without `-p` the program considers `RATPOINTS_DEFAULT_NUM_PRIMES` of them to begin with and looks further only for a curve that needs it (Section 8.1), whereas with `-p` it looks at exactly M . The program will, if necessary, reduce M , N , and n (in this order) to ensure that $n \leq N \leq M \leq \text{RATPOINTS_NUM_PRIMES}$. The latter is 53 by default (the number of odd primes $< 2^8$), and 30 or 96 if `PRIME_SIZE` is set to 7 or 9, etc.

M has a price in memory, because the sieving tables for a prime p take p rows of p bits and so grow with p^2 : the primes below 128 need about 5 MB, those below 1024 about 1.7 GB (twice that with 512-bit registers). Only the primes that M actually asks for are reserved, so raising `PRIME_SIZE` in order to have larger primes available costs nothing as long as M is left alone.

If `-n` or `-N` is left out, the corresponding number is chosen from the curve, as described in Section 8.1. The two constants that govern that choice can themselves be set:

`-r x -R e`

sets the target number of numerators surviving the first sieving stage per 64-bit word to x (a real number; this is what decides n), and the number of primes added to n to obtain N to e . The mnemonic is that `-r` and `-R` are to the automatic choice what `-n` and `-N` are to the explicit one. Both are machine-dependent, and giving them on the command line means they can be retuned without recompiling; the compiled-in defaults are `RATPOINTS_SURVIVORS_PER_WORD` and `RATPOINTS_SP2_EXTRA` in `ratpoints.h`.

Three further options control the same choice, and all three exist so that what they switch on can be switched off again and measured:

`-U u -C c -A a`

The first sets the length of run, in 64-bit words of numerators, at which a second-stage prime just pays for setting itself up; `-R` is then the number of extra primes for a run much longer than that, and a shorter run gets proportionately fewer (Section 8.1). `-U 0` makes `-R` a flat offset again, as it was before version 3.0.0. The second sets what building one row of a sieving table costs, in units of what one first-stage prime costs per word, which is what lets the program prefer a smaller

prime to a larger one of the same quality; `-C 0` ranks the primes by quality alone, as before. The third says how much of the choice the program should revise while it runs (Section 8.3): `a = 0` nothing, `a = 1`, the default, the number of primes in the third stage, and `a = 2` the number in the second stage as well. It revises nothing it did not choose in the first place, so `-n` or `-N` has the same effect as `-A 0`.

A third sieving stage runs after those two, on the numerators that survive them. It tests each of them against further primes one at a time, computing the position in the table of admissible residues rather than reading a bit out of a sieving table, so it needs no table and can use primes for which building one would be out of the question. How many primes it uses is chosen from the curve, and

`-P k`

overrides that with exactly `k` of them; `-P 0` switches the stage off, and

`-Q q`

sets instead the machine constant the choice turns on, namely what carrying one more prime through the denominators costs as a fraction of one exact test. It is to `-P` what `-r` is to `-n`. The choice weighs what a prime removes, which is a fraction of the survivors, against what it costs, which is a little work per survivor and a little more per denominator, so it depends on how many survivors a denominator has left by then — many for a curve with many rational points, few for a random one, and more at a large height bound than at a small one. On a random curve at the default height bound the stage does nothing at all, which is as it should be. Section 8.2 has the rule; the two machine constants in it are `RATPOINTS_SP3_PER_SURVIVOR` and `RATPOINTS_SP3_PER_DENOM` in `ratpoints.h`.

Both of those constants are fractions of what one exact test costs, and that is not the same number for every curve: the test evaluates a binary form of the given degree at one pair of integers and takes an integer square root, so it costs more at a high degree, and more again when the coefficients are large enough to push the arithmetic into another limb. The program estimates it from the degree, the largest coefficient and the height bound, all of which it knows before it looks at a single prime, and

`-W w`

overrides the estimate with `w` rdtsc cycles. `-W 306` is what versions before 3.0.0 did in effect, namely assume that every curve costs what a degree-6 curve with small coefficients costs; it is the way to measure what the estimate is worth. Section 8.2 has the formula, and `-v` prints what it gives for the curve in hand.

There are two more parameters that can be set.

`-F D`

sets the maximal number of ‘forbidden divisors’. If the degree is even, the denominator of any x -coordinate of a rational point is restricted at certain primes p : it cannot be divisible by p when the leading coefficient is a non-square mod p , and when p divides the leading coefficient its valuation at p is restricted instead — the program works out which valuations can occur at all (Section 7.5.3), which typically rules out a power of p , or one particular valuation, or p itself. Each such prime counts as one forbidden divisor. The program constructs a list of them against which to check the denominators, taking the primes up to the square root of the height bound, and this option caps the length of the list; the default is 64, which the list reaches at a height bound of about half a million. A prime in the list costs a few instructions per 64 denominators. With `-v` it prints what it found.

The other option is

`-S S` or `-s .`

In the first form, it sets the number of refinement (interval halving) steps in the isolation of the

real components to S . This part of the program computes a Sturm sequence for the polynomial and uses it in order to find a union of intervals that contains the intervals of positivity of the polynomial. This is done by a successive subdivision of the interval $]-\infty, +\infty[$. The number S sets the recursion depth. S can be omitted; then it is given a default value. The option `-s` skips the Sturm sequence computation completely. This also has the effect of removing most of the check for squarefreeness.

5.6 Switching off optimizations

The options `-k` and `-j` can be used to prevent the program from reversing the polynomial, which it usually does if this will lead to faster operation (`-k`), and to prevent the program from using the Jacobi symbol test on the denominators (`-j`). This last test extends the ‘forbidden divisors’ method described in Subsection 5.5 above by computing the Jacobi symbol (l/d) , where l is the leading coefficient and d is the odd and coprime-to- l part of the denominator. If the symbol is -1 , the denominator need not be considered. The use of these options may be questionable, unless you want to see how much performance is gained by using these optimizations.

5.7 Switching off exact testing of points

The option `-x` will prevent the program from checking the potential points that survive the sieve whether they really give rise to rational points. This implies that the points that are output may not actually be points. It also effectively sets the `-y` option, since the y -coordinates are not computed.

5.8 Overriding previous options

The options `-I`, `-Y`, `-Z`, `-S`, `-K`, `-J`, `-X` can be used to cancel the effect of the corresponding lower-case option (and thereby restore the default behavior), when this occurs earlier in the list of options. This may be useful if you want to set a default behavior that is different from what the program does out-of-the-box (e.g., in a shell script), but want the caller to be able to override this change.

6 How to use the library

It is possible to use the `ratpoints` machinery from within your own programs.

The library `libratpoints.a` provides the following functions.

```
long find_points(ratpoints_args*,
                int proc(long, long, const mpz_t, void*, int*),
                void*);

void find_points_init(ratpoints_args*);

long find_points_work(ratpoints_args*,
                    int proc(long, long, const mpz_t, void*, int*),
                    void*);

void find_points_clear(ratpoints_args*);
```

The passing of arguments to these functions is via the `ratpoints_args` structure, which is defined as follows.

```

typedef struct { mpz_t *cof; long degree; long height;
                 ratpoints_interval *domain; long num_inter;
                 long b_low; long b_high; long sp1; long sp2; long sp3;
                 double survivors_per_word; long sp2_extra; double sp2_u0;
                 double cost_table; long adapt;
                 long sp3_extra; double sp3_per_denom; double check_cost;
                 long array_size;
                 long sturm; long num_primes; long max_forbidden;
                 unsigned int flags;
                 long sp1_used; long sp2_used; long sp3_used; ...}
    ratpoints_args;

```

The dots at the end stand for additional fields that are used internally and are not of interest here. All fields up to and including `flags` are inputs, and they come back from a call as they went in, with three exceptions that are described below: the coefficient array and degree, which on return describe the polynomial the program worked with; `num_inter` and `domain`, which on return describe the region it searched; and the flag bits that report on the run. (Up to version 2.2.4 the program also wrote the values it chose or normalised into the other input fields, so that a caller who fills the structure once and then loops over curves had to reset them before every call; this is no longer necessary.) Every input field has to be set before the call; a negative value asks for the compiled-in default, or for the value to be chosen from the curve, as described below. The three `_used` fields are outputs, see below.

When calling `find_points`, the `cof` field must point to an array of size `degree + 1` of properly initialized gmp integers; these are the coefficients of the polynomial. The program can alter the values of the integers in this array: it may reverse the polynomial, and it drops a zero constant term on that occasion, in which case `degree` is lowered by one so that `cof` and `degree` still describe the polynomial the program worked with (the `RATPOINTS_REVERSED` flag says when this has happened). Leading zero coefficients are dropped in the same way, lowering `degree`, without a flag. To prevent the reversal, set the `RATPOINTS_NO_REVERSE` bit in `flags`; this may result in a loss of performance, though.

`height` gives the height bound. It is an error for the degree or the height bound to be nonpositive; in this case the function returns the value `RATPOINTS_BAD_ARGS`.

The field `domain` must contain a pointer to an array of `ratpoints_interval` structures of length at least `num_inter` plus `degree`. This array gives (in its first `num_inter` entries) the intervals for the search region. The type `ratpoints_interval` is just

```
typedef struct {double low; double up;} ratpoints_interval;
```

the meaning of this should be clear. In `ratpoints`, this is set via the `-l` and `-u` options. Usually, you do not want to restrict the range of x -coordinates; then you set `num_inter = 0` (but you still have to fill `domain` with a valid pointer to at least `degree` intervals; a null pointer makes the call return `RATPOINTS_BAD_ARGS`, since the program stores the default interval `]-height,height[` there. If `sturm` is negative, one interval is enough.) The program intersects the given intervals with the region where f is positive (unless `sturm` has a negative value, which may lead to a performance loss) and searches the result; on return, `num_inter` and the first `num_inter` entries of the array describe that result, which is why the array has to be longer than the input. So `num_inter` (and `domain`, if you restrict the search) must be set before every call.

`b_low` and `b_high` carry the lower and upper bounds for the denominator; if non-positive, they are set to 1 and the height bound, respectively. In `ratpoints`, these fields are set by the `-dl` and `-du` options.

sp1 and sp2 specify the number of primes to be used in the first sieving stage and in both sieving stages together, respectively. If negative, they are chosen from the curve, as described in Section 8.1; this is what the command line program does when -n and -N are absent. In ratpoints, these fields are set by the -n and -N options. What was actually used, chosen or not, is reported in sp1_used and sp2_used. The fields survivors_per_word and sp2_extra hold the two constants that govern that choice (options -r and -R); if negative, the values compiled into ratpoints.h are used. So do sp2_u0 and cost_table (options -U and -C), which say how long a run has to be for a second-stage prime to pay for setting itself up and what one row of a sieving table costs; zero in either switches off the corresponding refinement. adapt set to zero stops the program revising these numbers as it runs (option -A, Section 8.3); it revises them only when it chose them, so a non-negative sp1 or sp2 has the same effect.

sp3_used reports the number of primes used in all three stages together, so that sp3_used - sp2_used of them belong to the third. (sp3 itself is a working field of no interest to the caller.) What goes in is sp3_extra, which fixes that difference when it is non-negative (option -P) and otherwise leaves the number to be chosen from the curve, and sp3_per_denom (option -Q), which is the machine constant that choice turns on and takes the value compiled into ratpoints.h when negative. check_cost (option -W) says what one exact test costs, in rdtsc cycles, and both of those constants are fractions of it; when negative, which is the usual case, the program estimates it from the degree, the coefficients and the height bound as described in Section 8.2. Similar statements are true for num_primes (option -p), sturm (options -s, -S) and max_forbidden (option -F). The field array_size specifies the maximal size (in bit arrays; a bit array contains 64, 128, 256 or 512 bits, depending on the configuration) of the array that is used in the first sieving stage. If non-positive, it is set to a default value. The various default values are defined at the beginning of the header file ratpoints.h, where also the maximal degree of the polynomial is set.

The flags field holds a number of bit flags.

- RATPOINTS_NO_CHECK — when set, do not check whether the surviving x-coordinates give rise to rational points (set by the -x option to ratpoints).
- RATPOINTS_NO_Y — only list x-coordinates (in the form $(x : z) \in \mathbb{P}^1(\mathbb{Q})$) instead of actual points (with a y-coordinate); this is set by the -y option to ratpoints.
- RATPOINTS_NO_REVERSE — when set, do not allow reversal of the polynomial (set by the -k option to ratpoints).
- RATPOINTS_NO_JACOBI — when set, prevent the use of the Jacobi symbol test (set by the -j option to ratpoints).
- RATPOINTS_VERBOSE — when set, causes the procedure to print some output on what it is doing (set by the -v option to ratpoints).

There are some other flags that are used internally. One of them might be of interest:

- RATPOINTS_REVERSED — when set after the function call, this indicates that the polynomial has been reversed (and the contents of the cof array have been modified).

The main vehicle for passing information back to the caller is the proc function argument together with the pointer info. This function

```
int proc(long x, long z, const mpz_t y, void *info, int *quit)
```

is called whenever a point was found. x, y and z are the coordinates of the point (where y is a gmp integer). info is the pointer that was passed to find_points; this can be used to store information that should persist between calls to proc. If *quit is set to a non-zero value, this indicates that find_points should abort the point search and return immediately; otherwise the search continues. This is how the -1 option works. The return value is taken as a weight for counting the points; usually it will be 1.

The usual framework for using `find_points` is as follows.

```
(...)
#include "ratpoints.h"

(...)

mpz_t c[RATPOINTS_MAX_DEGREE+1]; /* The coefficients of f */
ratpoints_interval domain[2*RATPOINTS_MAX_DEGREE];
        /* This contains the intervals representing the
           search region */

/*****
 * function that processes the points
 *****/

typedef struct {...} data;

int process(long x, long z, const mpz_t y, void *info0, int *quit)
{ data *info = (data *)info0;

    (...)

    return(1);
}

/*****
 * main
 *****/

int main(int argc, char *argv[])
{
    long total, n;
    ratpoints_args args;

    long degree      = 6;
    long height      = 16383;
    long sieve_primes1 = RATPOINTS_DEFAULT_SP1;
    long sieve_primes2 = RATPOINTS_DEFAULT_SP2;
    long num_primes   = RATPOINTS_DEFAULT_NUM_PRIMES;
    long max_forbidden = RATPOINTS_DEFAULT_MAX_FORBIDDEN;
    long b_low        = 1;
    long b_high       = height;
    long sturm_iter    = RATPOINTS_DEFAULT_STURM;
    long array_size    = RATPOINTS_ARRAY_SIZE;
    int no_check       = 0;
    int no_y           = 0;
    int no_reverse     = 0;
    int no_jacobi      = 0;
    int no_output      = 0;
```

```

unsigned int flags = 0;

data *info = malloc(sizeof(data));

/* initialize multi-precision integer variables */
for(n = 0; n <= degree; n++) { mpz_init(c[n]); }

(...)

{ /* set up polynomial */
    long k;
    for(k = 0; k < 7; k++) { mpz_set_si(c[k], ...); }

    args.cof          = &c[0];
    args.degree       = 6;
    args.height       = height;
    args.domain       = &domain[0];
    args.num_inter    = 0;
    args.b_low        = b_low;
    args.b_high       = b_high;
    args.sp1          = sieve_primes1;
    args.sp2          = sieve_primes2;
    args.survivors_per_word = -1.0; /* negative: the compiled-in default */
    args.sp2_extra     = -1;
    args.sp2_u0        = -1.0;
    args.cost_table    = -1.0;
    args.adapt         = -1;
    args.sp3_extra     = -1;          /* negative: chosen from the curve */
    args.sp3_per_denom = -1.0;
    args.check_cost    = -1.0;
    args.array_size    = array_size;
    args.sturm         = sturm_iter;
    args.num_primes    = num_primes;
    args.max_forbidden = max_forbidden;
    args.flags         = flags;

    info->... = ...;
    (...)

    total = find_points(&args, process, (void *)info);
    if(total == RATPOINTS_NON_SQUAREFREE)
    { ... }
    if(total == RATPOINTS_BAD_ARGS)
    { ... }

    (...)
}

```

```

/* clean up multi-precision integer variables */
for(n = 0; n <= degree; n++) {mpz_clear(c[n]); }

return(0);
}

```

If points are to be searched on many curves of the same degree, then it is slightly more efficient to use the sequence

```

args.degree = degree; /* this information is needed */
find_points_init(&args);

for( ... )
{ ...
  total = find_points_work(&args, process, (void *)info);
  ...
}

find_points_clear(&args);

```

This avoids the repeated allocation and freeing of memory. Inside the loop only `cof`, `degree` and `num_inter` (and whatever else changes from curve to curve) need to be set; the other input fields keep their values across calls, since version 3.0.0. `rptest.c` with the `-O` option does exactly this, and make `test1once` checks that the points found are the same as with the fields set once per curve.

For practical examples, see `main.c` (the code that wraps `find_points` for the command line program `ratpoints`) or `rptest.c` (which runs `find_points` on some test data).

7 Implementation

7.1 Overview

Let $F(x, z)$ be the binary form of even degree corresponding to the polynomial on the right hand side of the curve equation. The basic idea is to let b run from 1 to the height bound H , and for each b , let a run from $-H$ to H , and finally, for each coprime pair (a, b) check if $F(a, b)$ is a square.

Of course, in this form, this would take a very long time. To speed up the process, we try to eliminate quickly as many pairs (a, b) as possible before the actual test. This can be done by ‘quadratic sieving’ modulo several primes: if $F(a, b)$ is a square, it certainly has to be a square mod p , and so we can rule out all (a, b) that do not satisfy this condition. Another ingredient is to represent (for a fixed b) the various values of a by bits and treat all the bits in a ‘bit array’ (of 64 bits, or of 128, 256 or 512 bits when SSE/AVX instructions are used) in parallel. For this, we organize the sieving information for each prime p into p arrays of p words each, one such array for every $b \bmod p$, such that the j th bit in this array is set if and only if $F(j, b)$ is a square mod p .

For each ‘denominator’ b , we then set up an array of words whose bits represent the range $-H \leq a \leq H$ (aligned so that $a = 0$ corresponds to the 0th bit of a word); the bits are initially set. Then for each of the sieving primes p , we perform a bit-wise *and* operation between this array and the sieving information at p . In a first stage, this is done on the whole array; after this first stage, each remaining (‘surviving’) non-zero word in the array is subject to tests with more primes. If some bits are still set after this second stage of sieving, the corresponding pairs (a, b) (if coprime) are

tested against a few further primes one at a time, without sieving tables (the third stage, since version 3.0.0), and those that pass are then checked exactly.

In the following subsections, we discuss a number of improvements that were made.

7.2 Sorting the primes

This idea is due to **Colin Stahlke**. We do not just take the first so many primes in increasing order for the sieving, but we first compute the number of points the curve has mod p for a number of primes p and then sort the primes according to the fraction of x -coordinates that give points. We then take those primes for the sieving that have the smallest fraction of ‘surviving’ x -coordinates. In this way, we need fewer sieving primes to achieve a comparable reduction of point tests. Since version 3.0.0 the candidates are ranked by what they say per unit of what they cost rather than by the fraction alone (Section 8.1), and they are not all primes: the odd prime powers and the products of small primes up to 64 take part, with the product of their factors’ fractions — a table row only has to be periodic in the bit index with the modulus as its period, and the row of a product of coprime moduli is the *and* of its factors’ rows — and a candidate sharing a prime with one already taken is passed over.

7.3 Using connected components

We note that we can only have points when $F(a, b)$ is non-negative. If we can determine intervals on which $f(x) = F(x, 1)$ is negative, then we do not have to look for points in these intervals. The necessary computations can be performed exactly, by computing a Sturm sequence for f and counting the number of sign changes at various points, see [Coh, Thm. 4.1.10]. This can in particular tell us whether F is negative definite, in which case we have already proved that there are no rational points in the curve. If there are real zeros, we use a subdivision method in order to find a collection of intervals containing the projections of the connected components of the curve over $\mathbb{R}$.

7.4 Using 2-adic information

John Cremona suggested that in some cases, one can determine beforehand that all points will have odd ‘numerators’ a , and so we can pack the bits more tightly by only representing odd numbers. This approach can be extended. We first find all solutions mod 64 (mod 16 before version 3.0.0). Then for every residue class mod 64 of the denominator b , we can find the residue classes mod 64 of potential numerators a . If there are none, then we can eliminate b as a denominator altogether. If all potential a ’s lie in one residue class modulo a power of two — they are all odd for an even denominator, and on many curves those of the odd denominators lie in one class modulo 4 or 8 — we can restrict the sieving to such numerators: a bit of the bit arrays then stands for every 2^k -th integer, for the largest such k (2, 4 and 8 occur on the test curves), and the arrays are that much shorter. Before version 3.0.0 only the restriction to one parity was made.

7.5 Elimination of denominators

Depending on the equation, certain denominators can be excluded. We have seen an instance of this in the previous subsection, but we can also work with odd primes.

7.5.1 Odd degree and $\pm$ monic

If the polynomial has odd degree and leading coefficient ± 1 , then the denominator has to be a square. This reduces the time complexity tremendously.

7.5.2 Odd degree general

In general, when f has odd degree, the denominator has to be ‘almost’ a square: it must be a square times a (squarefree) divisor of the leading coefficient, and there are further restrictions on the parity of the valuation at p , for p dividing the leading coefficient, when this valuation is sufficiently large. This gives the same type of time complexity as in the monic case, but with a larger constant.

7.5.3 Even degree

Let l be the leading coefficient. If l is not a square, we can exclude denominators divisible by a prime p such that l is a non-square mod p : for $p \mid b$ we have $F(a, b) \equiv la^n \pmod{p}$ with $p \nmid a$, which is a non-square as well. In addition, we use the necessary condition that the Jacobi symbol $(\frac{l}{b'})$ must be $+1$, where b' is the part of b coprime to $2l$. Since version 3.0.0 that symbol is not computed by a binary algorithm for each denominator but as a product of Legendre symbols: for $l = \pm 2^v \prod_i q_i^{e_i}$, with $\epsilon = 1$ if $l < 0$ and 0 otherwise, it is $(\frac{-1}{b'})^\epsilon (\frac{2}{b'})^v \prod_i (\frac{q_i}{b'})^{e_i}$, where by quadratic reciprocity each factor is the Legendre symbol of $b' \pmod{q_i}$ modulo q_i times a sign that depends on $b' \pmod{4}$, and the first two factors depend on $b' \pmod{8}$. The program tabulates the non-squares modulo each q_i once per curve, so that the test costs one multiplication and one table look-up per prime factor of l ; when l has a prime factor above 1021, or more odd prime factors than the tables allow for (sixteen, or 8 KB), or the denominators exceed 2^{32} , it falls back on the binary algorithm.

When p divides l , the congruence says nothing, since $F(a, b)$ is then $0 \pmod{p}$; the valuation does. Write $v_p(b) = m \geq 1$ and let w_j be the valuation at p of the coefficient of b^j in F , so that the term with b^j has valuation $w_j + jm$. If one term has strictly smaller valuation than all the others, it determines $v_p(F(a, b))$ and the residue of $F(a, b)/p^v \pmod{p}$. If that valuation is odd, no denominator with $v_p(b) = m$ can occur. If it is even and j is even, then a^{n-j} and $(b/p^m)^j$ are both squares mod p , so none can occur when the unit part of that coefficient is a non-square mod p . (For odd j the factor $(b/p^m)^j$ takes non-square values as well as square ones, and nothing follows.) When two terms tie for the smallest valuation they may cancel, and nothing follows either; indeed that is what happens at the points of such a curve. For large m the leading term wins by itself, so when $v_p(l)$ is odd, or even with a non-square unit part, every sufficiently large valuation is excluded: the denominator cannot be divisible by a certain power of p . The commonest cases are $v_p(l) = 1$ with $p \mid l_{n-1}$, which excludes p itself, and $v_p(l) = 1$ with $p \nmid l_{n-1}$, where the two top terms tie at $m = 1$ and p^2 is excluded. Below such a tail there may be isolated excluded valuations, and there may be some when l is a square: $l = 9u$ with $3 \nmid l_{n-1}$ excludes $v_3(b) = 1$ whatever u is.

While computing the sieving information, we make a list of these primes, each with the set of valuations it excludes, which we then use to eliminate denominators: a prime that is excluded altogether is tested for with bit arrays in the manner of the sieve, the others by computing the valuation of the denominator. The bit arrays cover every excluded prime up to the square root of the height bound (or up to 1021), building the patterns for the primes beyond the compiled sieving primes as needed: past the square root a denominator can have at most one excluded prime left, and the Jacobi symbol sees that one, whereas it lets an even number of them through.

Up to version 2.2.3 the valuation test never ran. It sat behind a guard that asked for l to be a non-square modulo p , and $l \equiv 0$ is a square. That was fortunate, because it was also wrong: in working out from which power of p on the leading term decides the valuation, it left out every coefficient that is a unit at p , which are the ones that decide it.

7.6 Reversing the polynomial

The exclusion of denominators described above is more or less effective, depending on the situation; it is the better the earlier the case was described. Since the height of the points does not change under the transformation $(x, y) \rightarrow (1/x, y/x^{\lceil n/2 \rceil})$, we can as well search on the reversed polynomial $F(z, x)$. This will result in a speed-up if the reversed polynomial belongs to a ‘better’ class than the original.

7.7 Some general remarks

Modern processors are very fast when doing basic things like moving data around between registers or simple arithmetic operations (addition, subtraction, shift, . . . , comparison). Even memory access can be fast when there is no cache miss. On the other hand, integer division is a slow process. It turned out that quite some improvement of the performance was possible by removing as many instances of integer division operations as reasonably possible.

For example, the gcd that tests a candidate for lowest terms is a binary gcd without a data-dependent branch; the Jacobi symbol test on the denominators is evaluated as a product of Legendre symbols read from tables built once per curve (Section 7.5.3); the residue classes of b modulo the various sieving primes are computed by a multiplication with a precomputed reciprocal rather than by a division; and most of the testing of ‘forbidden divisors’ of b is done with bit arrays similar to those used for sieving the numerators, a word of 64 denominators at a time.

8 Fine-tuning the parameters

Most of what the sieve does is decided from the curve before a single numerator is sieved: which moduli to sieve with and how many in each of the three stages, how long the run is going to be, and what one exact test will cost. The rules weigh running-time costs against each other, so they rest on a few constants that describe the machine rather than the curve. Reasonable values are compiled into `ratpoints.h`, `make tune` measures better ones (Section 8.4), and each of them can be set on the command line for a single run. This section says what the rules are and what the constants mean; Section 7 says how the sieve works.

8.1 How the moduli of the first two stages are chosen

While collecting the sieving information, the program computes for each prime p — and for each odd prime power and product of small primes up to 64, which sieve in the same way (Section 7) — the density r_p , the proportion of residues a modulo p for which $f(a)$ is a square. The expected proportion of numerators surviving a set of moduli is the product of their densities, so a modulus with a small density says more.

The moduli are not ranked by density alone, because they do not cost the same. A modulus m needs a sieving table of m rows, built once for each class of denominators modulo m that turns up — at most m times, however long the run — and an entry per denominator in the list of reduced denominators; per word of numerators both fall like $1/U$, where U is the number of words the run is going to sweep. The program computes U and the number of denominators from the search region before it looks at a modulus (the numerators of one denominator are a piecewise linear function of it, so a sample over the denominators is accurate to a fraction of a per cent), and ranks the moduli by

$$\frac{\text{cost}(m)}{-\log r_m},$$

the cost per word divided by the information. At equal quality a smaller modulus is preferred, and at a small height bound strongly so. `RATPOINTS_COST_TABLE` (option `-C`) is what building one row of a table costs, in units of one first-stage AND per word; `-C 0` ranks by density alone.

The first stage is filled from the top of that ranking for as long as a modulus pays: it is taken while the survivors per word it removes, weighted by what a survivor costs in the stages after it, exceed `RATPOINTS_SURVIVORS_PER_WORD` (option `-r`) times what the modulus costs per word — one AND, plus its tables, its per-denominator entries and the fetch of its row, spread over the run. That threshold is the ratio of two machine costs and the most important constant of all. With it the first stage stops at a few thousandths of a surviving bit per word whatever the curve: at about seven moduli for a run of a few hundred numerators per denominator, at about ten at the height bound 16383 of the test suites. The second stage then gets

$$sp2 = sp1 + \frac{RATPOINTS_SP2_EXTRA}{1 + RATPOINTS_SP2_U0/U}$$

further primes, capped by the number that carry any information at all. `RATPOINTS_SP2_EXTRA` (`-R`) is the number an arbitrarily long run wants, and `RATPOINTS_SP2_U0` (`-U`) is the length of run at which setting a prime up costs as much as sieving with it, so that a short run gets proportionately fewer; `-U 0` makes the offset flat. The second stage keeps to primes and has a ranking of its own, since a prime there is applied only to the bit arrays that survived the first stage, and its table therefore weighs some three times more heavily against it.

A curve with very many rational points makes f a square modulo every residue for the smallest primes. Those primes say nothing and are discarded, and the thirty primes the program looks at by default may not leave enough for the rule to have what it wants; the program then looks at further primes, one at a time, until they suffice or the table ends. Every prime added can only lower `sp1` and raise the number of informative primes, so this stops of its own accord, and no random curve ever triggers it. Giving `-p` explicitly switches it off: an explicit number of primes is a hard limit, as are `-n` and `-N` for the two stages.

8.2 How many primes the third stage uses

The third stage tests the survivors of the first two one at a time, against further primes and without a sieving table. A prime saves the fraction $1 - r_p$ of the survivors that reach it, each of which would otherwise go to the exact test with its multi-precision arithmetic and its integer square root; it costs a multiplication and a reduction for every survivor tested, and one step per denominator to keep the denominator reduced modulo it, whether or not a survivor turns up. So a prime is worth adding while

$$S(1 - q_s - r_p) > q_d,$$

where S is the expected number of survivors a denominator still has and q_s , q_d are the two costs as fractions of one exact test — `RATPOINTS_SP3_PER_SURVIVOR` and `RATPOINTS_SP3_PER_DENOM` in `ratpoints.h`; option `-Q` sets the second. Since S falls by the factor r_p with each prime, this stops of its own accord. On a random curve at a small height bound the stage uses no prime at all; on a curve with many rational points at a large one it uses twenty or more, and looks further up the table of primes whenever none in hand pays but a prime of the density the curve has been offering would. `-P k` fixes the number at k instead, and `-P 0` switches the stage off.

S is estimated from the number of numerators a denominator has to consider, thinned by the 2-adic information, by the moduli of the first two stages, and by the proportion of numerators coprime to the denominator — that last because the test for common factors runs before this stage, and has to: a numerator sharing a factor with the denominator stands for a fraction a smaller denominator has dealt with already, so f takes the same value there and every prime accepts it.

What one exact test costs depends on the curve, not only on the machine. The test evaluates the binary form $F(a, b)$ of degree d by Horner's rule and takes an integer square root, so it costs a step per degree on numbers that grow to about $\log_2 |c| + d \log_2 H$ bits, plus a fixed amount for every further limb of the root; over the range from degree 3 with small coefficients to degree 14 with 400-bit ones it varies by a factor of ten. The program estimates it from the degree, the largest coefficient and the height bound before it looks at a prime — the constants are `RATPOINTS_CHECK_STEP`, `_LIMB`, `_CALL` and `_ROOT`, relative to `RATPOINTS_CHECK_REFERENCE`, the value for a degree-6 curve with small coefficients — and divides the two fractions above by the ratio, so that a curve whose test is dearer gets more primes. Only ratios enter, so a machine on which every test costs more needs no change; `-W w` overrides the estimate with w cycles, and `-v` prints what the estimate gives for the curve in hand.

8.3 Correcting the choice while the program runs

The rules above use the densities alone, and they miss one thing by construction: if a/b is a rational point, so is ka/kb for every k — it is the same rational number, with the same value of f — and it passes every prime test there is. A floor of survivors therefore outlives any amount of sieving, and only the test for common factors removes it. The program can see the floor because it is sieving: it counts the bit arrays that survive the first stage, the survivors of the second and the candidates that pass the test for common factors, all of them on paths that are taken rarely, and fits

$$S(n) = \text{floor} + \text{chance} \cdot R(n)$$

to the counts, with $R(n)$ the product of the n smallest densities. It revises its choice after the first million words of numerators and then whenever twice as many have been swept. By default (`-A 1`) it revises only the third stage, where the number of survivors a denominator brings becomes a count rather than an estimate; `-A 2` reconsiders `sp2` as well, and `-A 0` switches the revision off. Changing the number of primes part-way through a run is safe: sieving only ever removes numerators that cannot be points, so what is found does not depend on it. Nothing that was fixed with `-n`, `-N` or `-P` is revised.

8.4 Retuning for your machine

The five constants — the threshold, the offset, the run length at which a second-stage prime pays for setting itself up, what building a table row costs, and the third stage's cost per denominator — were fitted on the author's laptop by minimising the sum of the times of `make test1` and `make test1many`, which between them cover the two regimes that occur in practice, random curves and curves with many rational points. To repeat that on your machine, run

```
> make tune
```

It sweeps the five constants, in that order, over both test sets and writes the values it finds to `tuning.mk`, which the `Makefile` includes, so that a subsequent `make all` builds with them. No source file is touched, and deleting `tuning.mk` restores the values compiled into `ratpoints.h`. The file records the configuration it was measured for and is ignored, with a warning, if you later build with another register width or `PRIME_SIZE`.

It is deliberately not part of `make all`: it takes several minutes (about three for the 256-bit build on the author's machine, four for 128 bits, six for the plain 64-bit one and eight for the 512-bit code emulated through narrower operations), wants an otherwise idle machine, and must not be run under `make -j`. Timing this reliably is the hard part. The cost surface is flat — anything within a factor of two of a good threshold costs under 3% — whereas a laptop under sustained load slows by a quarter as it warms up, so a straight sweep would measure the clock rather than

the settings. Each candidate is therefore timed back to back with the current settings and only the ratio is kept, the median over several rounds is used, and the current settings are themselves among the candidates, so that the ratio they measure against themselves shows how much noise is left. If that is more than 2% the run is declared inconclusive, and nothing is written unless the winner is better by more than 3%: a noisy run leaves a good default alone. Two consequences are worth knowing in advance. A run that reports that nothing beat the current settings has *succeeded*; on a build that is already tuned that is the right answer. And a configuration whose tests take long enough for the machine to throttle within a single round may report itself inconclusive however often it is run; raising `ROUNDS`, or using a machine that does not throttle, is then the remedy. The environment variables `ROUNDS` (how many times each setting is measured, default 3), `WARMUP` (seconds of load before anything is counted, 20), `MARGIN` (0.97) and `NOISE` (0.02) adjust this, and `R_VALUES`, `E_VALUES`, `U_VALUES` and `C_VALUES` give the ladders of candidates.

```
> make tunehigh
```

measures on `make testhigh` and `make testhighmany` instead, at a height bound of $2 \cdot 10^5$, where the two sieving phases are nine tenths of the run; at the 16383 of `test1` two fifths of the time on the random curves goes into work the constants have little or no effect on. Use it if the runs that matter to you are long ones. One timing there is two minutes rather than three seconds, so it does not sweep the whole ladder: it starts from the settings in force — which is to say from what `make tune` found, since both read and write the same `tuning.mk` — and asks only whether a factor of two either way in the threshold, in the run length or in the table cost, or two more or fewer second-stage primes, is better (`R_FACTORS`, `E_DELTAS`, `U_FACTORS` and `C_FACTORS` set that neighbourhood). Expect a couple of hours at the default three rounds, and run it again if it moved a value. The intended order is `make tune` first and `make tunehigh` afterwards; only a pair of tunings at run lengths a hundred times apart can separate the offset from the run length it is scaled by.

Should you prefer to do it by hand, the constants can be given for a single run with `-r`, `-R`, `-U`, `-C` and `-Q`, which `ratpoints` and `rpctest` both accept, and hard-coded by changing the definitions in `ratpoints.h` (followed by `make distclean all test`). `rpctest -z -T` prints the CPU time of a run and nothing else, which is what `tune.sh` uses; `-h H` changes its height bound, `-m m` repeats the computation `m` times, and `testdata.h` can be replaced by a collection of your own curves. Programs using the library set the same constants through the fields of `ratpoints.args` (Section 6).

9 Change log

Version 2.0 was released January 9, 2008.

Version 2.0.1 was released July 7, 2008. It fixes a bug that prevented the `'-1'` option to work properly.

Version 2.1 was released March 9, 2009. It makes use of the SSE instructions, so that the sieve can work on 128 bits in parallel (instead of on 64 or 32). On my laptop (with an Intel Core2 processor), `make test` runs about 25% faster than before.

Version 2.1.1 was released April 14, 2009. It fixes a bug that in some cases prevented an early abort when `*quit` was set by the callback function. Thanks to **Robert Miller** for the bug report and the fix. In addition, it is now checked that `__WORDSIZE == 64` before SSE instructions are used (in `rp-private.h`), since the program then assumes that a bit array consists of two unsigned longs. In this context, `__SSE2__` has been replaced by a new macro `USE_SSE`. A further fix eliminates unnecessary copying of sieving information when SSE instructions are not used (introduced in

version 2.1). This was almost always harmless, but could have resulted in memory corruption in the extreme case that all the information had to be computed for all the primes.

Version 2.1.2 was released May 27, 2009. It fixes some memory leaks. Thanks to **Robert Miller**.

Version 2.1.3 was released September 21, 2009. The library function `find_points` should now work without any bound on the degree of the polynomial. The degree bound for the `ratpoints` executable is now set to 100 by default (specified by `RATPOINTS_MAX_DEGREE`, used to be 10).

A bug in `rptest.c` (pointed out to me by **Giovanni Mascellani** and **Randall Rathbun**) that could lead to a segmentation fault was fixed on March 10, 2011.

Version 2.2 was released January 9, 2022. The main change is that `ratpoints` can now also work with 256-bit (and, in principle, 512-bit) registers when the CPU has the relevant capabilities. Thanks to **Bill Allombert** for doing a first implementation of this. In addition, the code has been cleaned up to some extent (e.g., the variants according to register sizes are now completely dealt with through macros defined in `rp-private.h`), and some more comments have been added. There are also some further optimizations; for example, the first sieving phase now uses many of the SSE/AVX registers in parallel, and on some architectures, the code uses a faster implementation for the test whether a bit array is zero. On my current laptop, the new code (using 256-bit words) is about twice as fast as the old one (2.1.3 using 128-bit words).

Version 2.2.1 was released January 18, 2022. This fixes a small bug that was introduced in the previous version, which could lead to the program missing points at the upper end of the search interval for the numerators. Thanks to **Bill Allombert**.

Version 2.2.2 was released May 15, 2023. This fixes a small bug that led to the polynomials being reversed incorrectly when the degree is odd. Thanks to **Nicolas Mascot** and **Bill Allombert**.

Work on the code starting with version 2.2.3 was done with the help of **Claude Code** (by Anthropic), using various models (Opus 5, Fable 5, and Fable 5.1).

Version 2.2.3 was released September 6, 2026. It fixes the code path using 512-bit registers, which had never been exercised. Two problems were found. The first one would have made the program crash. The second problem was that the `TEST` macro, which is used in the innermost loop of the second phase, still used the generic fall-back version, leading to a degradation in performance; it now uses `vpstestmq`, in analogy with the 256-bit version.

Since I have no CPU with AVX512F available, this was tested by compiling with `-DUSE_AVX512` but without `-mavx512f`, so that `gcc` expresses the 64-byte vectors through 256-bit operations. This exercises the whole code path. What this cannot check is the genuine 512-bit instructions. In the meantime, **Drew Sutherland** has run the two variants against each other on an AMD Ryzen 9 9950X3D (Zen 5 architecture), for a curve with many rational points and a height bound of 10^6 ; the 512-bit version was faster by 13 to 14%.

Version 2.2.4 was released September 19, 2026. It fixes bugs that a review of the code and running the new test suite turned up:

- A polynomial of degree 1 — given as such, or reached from $y^2 = c_2x^2 + c_1x$ by the reversal, or from a zero leading coefficient — made the Sturm sequence code read past the end of its arrays and crash; this affects library users as well.
- When the positivity region of f did not meet the search domain, the program returned before it had looked at the points at infinity, so that `ratpoints` `'-1000000000 0 0 1'` 45 printed nothing where `-s` printed `(1 : 0 : 0)`, and after a reversal the lost point is an affine point of the given curve.
- When no residue class of the denominator modulo 16 admitted any numerator, the array of numerator patterns was read before it had been filled; by default the answer was still

right, but with `-j -F 0` the program printed nondeterministic points outside the height bound after a run several thousand times longer than necessary.

- Several input fields of `ratpoints_args` were used as scratch space — the number of forbidden divisors found was written into `max_forbidden`, the intersected search intervals into `num_inter`, normalised values into `sp1`, `sp2` and others — so that a program which fills the structure once and then calls `find_points_work` for curve after curve, as the section on the library suggests, had the forbidden-divisor test switched off after the first curve without forbidden divisors (a factor of two in running time) and the first curve’s positivity region intersected into the search domains of the others, which loses points. The program no longer writes either, and the manual now says exactly which fields it does set; `make test1once` runs the test curves with the structure filled once, and `make test3` runs one or more invocations of `ratpoints` per bug against a reference.
- The library tested whether the leading coefficient is a square before it checked the coefficient pointer for `NULL`, so that a caller passing `NULL` crashed instead of getting the error code back.
- An empty coefficient string made the program report “Bug no. 1 — please report!” instead of refusing the input.
- At the very top of the range of a long, three things went wrong: a search interval that reached the last few hundred numerators below 2^{63} was dropped without a word, two computations on the position of its last bit having overflowed; the loop over every denominator never ended when the denominator bound was $2^{63} - 1$, since the increment wraps around; and the loops over the square denominators formed b^2 and $d b^2$, which overflow within $2 \cdot 10^9$ of it.
- The test on the valuation of a denominator $d b^2$ at a prime dividing the leading coefficient applied `abs`, the `int` function, to a `long`, so that for denominators above 2^{31} it judged the low 32 bits only and could exclude a denominator that carries a point: a point with denominator $3 \cdot 59827^2$ was missed.
- Not a bug, but wasted time: the loops over the square denominators started at 1 whatever the lower bound `-d1` said; they now start at the first square in the range.
- Also fixed: four left shifts of negative values (undefined in C, harmless with gcc so far), two arrays that could be declared with length zero under `-n 0 -N 0`, an overflow in the enumeration of the divisors of the leading coefficient for denominator bounds above about $4 \cdot 10^{15}$, and two constants in `ratpoints.h` that could not be overridden from the compiler command line. The library is binary compatible with 2.2.3.

Version 3.0.0 was released September 19, 2026. Its theme is that the program works out for itself what it used to be told, and the sieve has been reworked from the set-up to the exact test. Measured against version 2.2.4 on the same curves and the same laptop (cycles pinned to one core, medians of three interleaved rounds; 2.2.4 at its own defaults, which sieve with the primes below 128 where this version uses those below 256; both find the same points on every suite):

suite	2.2.4, Gcycles	3.0.0	ratio	faster by
<code>make test1</code> (random curves, 16383)	8.45	2.89	0.34	2.9
<code>make test1many</code> (point-rich, 16383)	11.27	4.80	0.43	2.3
<code>make testhigh</code> (random, $2 \cdot 10^5$)	412.7	190.8	0.46	2.2
<code>make testhighmany</code> (point-rich, $2 \cdot 10^5$)	764.4	138.3	0.18	5.5
random curves at height 4000	3.06	0.54	0.18	5.6
random curves at height 1000	2.06	0.21	0.10	10.1
point-rich curves at height 1000	0.44	0.15	0.34	3.0
random curves at height 200	0.96	0.12	0.13	8.0

Instructions fell to between 0.16 and 0.47 of 2.2.4's and mispredicted branches to between 0.03 and 0.40: at the small height bounds 2.2.4 was bound by its branches, and cycles fall two to three times more than instructions there. The changes, in outline; the sections named have the details.

- *The sieving parameters are chosen from the curve.* How many moduli each of the first two stages uses, and which ones, follows from the densities of admissible residues, from the length of the run and from what each modulus costs, with the moduli ranked by cost per information rather than by information alone (Section 8.1); `-n` and `-N` still fix the numbers. When the small primes say nothing, as on a curve with very many rational points, the program looks at further primes of its own accord, and the default table holds the primes below 256 (`PRIME_SIZE` 8, formerly 7) to give it something to reach for. The rules rest on five machine constants, set with `-r`, `-R`, `-U`, `-C` and `-Q` and measured by `make tune` and `make tunehigh` (Section 8.4), which write them to `tuning.mk`.
- *A third sieving stage* tests the survivors of the two bit-array stages one at a time against further primes, without sieving tables, so that it can use primes far up the list; how many is decided per curve (Section 8.2, option `-P`) from an estimate of what one exact test costs on that curve (`-W`). While it runs, the program counts what the sieve is finding and revises the choice (Section 8.3, option `-A`).
- *Composite moduli.* Odd prime powers and products of small primes up to 64 sieve alongside the primes: modulo 9 carries twice what 3 does, and one AND with the row of $45 = 9 \cdot 5$ carries what two large primes do (`RATPOINTS_COMPOSITE_MAX` in `ratpoints.h`).
- *2-adic information modulo 64* instead of 16, the even denominators decided exactly through the reversed polynomial, and the bit arrays of each denominator class packed with the largest stride 2^k its numerators allow rather than by parity alone (Section 7).
- *Denominators.* The test at a prime dividing the leading coefficient, derived from the Newton polygon of F , runs for the first time — it had sat behind a guard it could never pass, and was wrong besides (Section 7.5.3); the forbidden-divisor bit arrays cover every excluded prime up to the square root of the height bound, and the default of `-F` is 64; the Jacobi symbol is a product of table-driven Legendre symbols; the tests on $b \bmod 64$ are applied to 64 denominators at a time; the residues of b come from precomputed reciprocals, and the gcd that tests a candidate for lowest terms is branchless.
- *The sieving loops.* The first stage ANDs the 2-adic pattern in with its first modulus instead of pre-filling the bit arrays, and sieves the bit arrays left over after the last whole chunk of sixteen in legs of 8, 4, 2 and 1 registers instead of padding the range; the second stage steps over the empty bit arrays into a sentinel, reaches its table rows through a precomputed reciprocal and extracts the survivors by counting trailing zeros. The sieving tables are built four times faster, and the memory reserved for them follows the number of primes in use.
- *Build and portability.* A stamp file makes a change of `CCFLAGS1` or `PRIME_SIZE` rebuild what depends on it; `CCFLAGS128` selects `USE_AVX128`, which needs only SSE2 and is faster than the `USE_SSE` variant; the code assumes a 64-bit long and refuses to compile without one, so version 2.2.4 is the last one for 32-bit machines.
- *Library.* `ratpoints_args` has gained fields, so programs using the library must be recompiled; the input fields come back as they went in, except `num_inter` and `domain`, which describe the region searched, and `sp1_used`, `sp2_used` and `sp3_used` report what was used (Section 6).
- *Tests.* `make test` runs, besides the random curves, the point-rich ones (`test1many`), a hundred curves of degrees 3, 4, 7 and 8 (`testdegrees`), the argument structure filled once for all curves (`test1once`), one invocation per bug fixed (`test3`), a suite that exercises every branch of the code against a reference checked by brute force (`test4`) and the library

interface (testapi), and it fails when a test fails; testhigh, testhighmany, test4configs and coverage are run separately (Section 4.3).

- *Bugs fixed.* All of those listed under version 2.2.4, and in addition: the estimate of the run length fell to its floor when no odd denominator admitted a numerator, which switched the second and third stages off for such curves; at a height bound of 2^{62} the third stage's bias overflowed and rejected every point; an interval endpoint of the positivity region was scaled by a shift that is undefined at large bisection depths (now ldexp); and the library refuses a negative degree.

References

- [Coh] H. Cohen, *A course in computational algebraic number theory*, Springer GTM **138**, second corr. printing, 1995.
[gmp] The GNU Multiple Precision Arithmetic Library, <http://gmplib.org/>
[rat] The ratpoints-3.0.0 package. Available at
<http://www.mathe2.uni-bayreuth.de/stoll/programs/ratpoints-3.0.0.tar.gz>

MATHEMATISCHES INSTITUT, UNIVERSITÄT BAYREUTH, 95440 BAYREUTH, GERMANY.

Email address: Michael.Stoll@uni-bayreuth.de